Morning Intel — Monday, Feb 16, 2026

TL;DR

- **OpenClaw creator steipete joins OpenAI** — biggest ecosystem news; monitor governance and fork risk for our daily-driver tool
- **Qwen 3.5 (397B MoE) dropping open-source** — another frontier-class local model; Qwen3-Coder-Next 80B fits 8 GB VRAM
- **Chrome DevTools official MCP server launched** — direct agent browser control from Google; strong candidate to replace our stealth-browse stack

Top 5 Highlights

1. steipete → OpenAI (★★★★★)

Peter Steinberger, OpenClaw creator, joins OpenAI. OpenClaw going open/independent. HN: 960 pts, X trending. We run on this tool — watch for governance changes, community forks (ClaudeClaw already exists).

2. Qwen 3.5 Released — 397B-A17B MoE (★★★★★)

Alibaba's biggest open model yet. 17B active params via MoE. Coding variant (80B) fits consumer hardware. Will be open-sourced. Another strong local option beyond Llama.

3. Chrome DevTools MCP (★★★★★)

Official Google MCP server gives coding agents full DevTools access — debugging, perf, network, DOM. Works with Claude Code, Cursor, Copilot out of the box.

4. NVIDIA commoditizing voice AI (★★★★)

PersonaPlex adds natural conversational overlap. Undercutting OpenAI on voice pricing. Changes cost equation for voice agents.

5. Opus 4.6 community reception overwhelmingly positive (★★★★★)

Dominating r/ClaudeAI. "Goated all-around model." We're already running it — validates the choice.

New Bookmarks

0 new bookmarks tonight. All 241 bookmarks already cataloged from previous runs. Nothing new saved since last scan.

Timeline Discoveries — 17 Posts Captured

No.	Author	Summary
1	@steipete	Joining OpenAI — OpenClaw going independent. 27.5K likes.
2	@damianplayer	Mark Cuban on selling AI agents to SMBs — actionable XPERIENCE framework
3	@aakashgupta	NVIDIA making every voice AI API a commodity
4	@heygurisingh	Google CodeWiki — paste repo URL → navigable wiki
5	@NoahEpstein_	Kimi model near Opus level at fraction of cost
6	@cgtwts	Kimi Claw: browser-based OpenClaw workspace, 5K+ skills
7	@openclaw	OpenClaw 2026.2.15: Telegram streaming, Discord Components v2
8	@robjama	Anthropic marketing team uses Claude Code with sub-agents + memory
9	@bloggersarvesh	Claude + Local SEO = blue collar millionaires — XPERIENCE thesis
10	@mfishbein	40 hrs of lead prospecting demolished with Claude Code + Browserbase
11	@getdelve	AI agents closed SOC 2 audit in 19 days. 8.6K likes.
12	@CoFoundersNik	steipete: first one-person unicorn. Weekend project → unicorn in 3 months.
13	@johann_sath	OpenClaw auto-backups to GitHub every 2 hours — resilience pattern

Timeline vibe: Dominated by steipete/OpenAI news. Heavy OpenClaw ecosystem activity. Strong "agents for SMBs" signal. The agent ecosystem is maturing fast and the SMB opportunity window is wide open.

GitHub Trending — 14 Repos

Repo	What It Does	Action
chrome-devtools-mcp	Official Chrome DevTools MCP — full browser control for agents	Integrate
gh-aw	GitHub agentic workflows in natural language markdown → Actions	Explore
rowboat	AI coworker with knowledge graph from email/meetings, MCP-extensible	Explore
RAG-Anything	All-in-one RAG for any document type	Watch
chatterbox	SoTA open-source TTS by Resemble AI	Explore
mistral-vibe	Minimal CLI coding agent by Mistral	Watch
get-shit-done	Meta-prompting + context engineering for Claude Code	Explore
claude-skills	66 specialized Claude Code skills collection	Explore
zvec	Alibaba's in-process vector DB — lightning fast, no server	Watch
PocketFlow	100-line LLM framework — agents building agents	Watch
Kronos	Foundation model for financial market language	Watch
pyrefly	Facebook's Rust-based Python type checker — faster than mypy	Watch

Trends: MCP becoming the standard agent-tool interface. AI coding agents maturing (Mistral Vibe, GSD, skills collections). Knowledge graphs over cold RAG. Vector DB proliferation.

News & Trends

Top Stories

1. steipete → OpenAI (HN 960pts + X trending)
OpenClaw creator joins OpenAI. OpenClaw going independent. Fork risk real — ClaudeClaw already emerging. ★★★★★

2. Qwen 3.5 Released (r/LocalLLaMA 350+ pts)
397B-A17B MoE, will be open-sourced. Qwen3-Coder-Next 80B runs on 8 GB VRAM. ★★★★★

3. Opus 4.6 Praised (r/ClaudeAI 744+ pts)
"Goated all-around model." Community overwhelmingly positive. ★★★★★

4. NVIDIA PersonaPlex (X trending 1,300 posts)
Natural conversational overlap for voice AI. ★★★★

5. MiniMax-2.5 Runs Locally (r/LocalLLaMA 413 pts)
Another local model option worth evaluating. ★★★★

6. AI Training Costs –40%/yr (Karpathy)
Continued cost deflation — good for indie builders. ★★★★

7. LLM Agent Cost Curve is Quadratic (HN)
Important analysis on agent economics — directly relevant to our architecture. ★★★★

8. Elon vs Anthropic Drama (r/ClaudeAI 1,619 pts)
Highest-voted post. Could affect Claude's positioning. ★★★★

9. Audio AI: Small Labs Winning (HN 187pts)
Smaller labs outperforming big tech in audio/voice. GPU access is the differentiator. ★★★★

Macro Trends

- Open-source model explosion: Qwen 3.5, MiniMax-2.5, Ling-2.5-1T all in one weekend
- Agent cost awareness growing — "vibe coding" hitting economic reality
- OpenClaw ecosystem at inflection point — creator leaves, community builds alternatives
- Voice/Audio AI heating up as next frontier
- AI training cost deflation becoming consensus (40%/yr)

Action Items

High Priority

- ☐ Integrate Chrome DevTools MCP into Claude Code config (GitHub: chrome-devtools-mcp)
- ☐ Update OpenClaw to 2026.2.15 — Telegram streaming + nested subagents
- ☐ Monitor OpenClaw governance/fork situation post-steipete departure

Medium Priority

- ☐ Evaluate Chatterbox TTS as potential claw-speak replacement
- ☐ Explore get-shit-done context engineering patterns for AGENTS.md
- ☐ Read the quadratic agent cost curve analysis (blog.exe.dev)
- ☐ Check out Mark Cuban SMB agent framework for XPERIENCE positioning
- ☐ Explore Rowboat knowledge graph patterns for memory system improvements

Low Priority / Watch

- ☐ Track Qwen 3.5 open-source release for local model evaluation

? Evaluate claude-skills collection for useful additions

? Try Google CodeWiki for project documentation

? Consider GitHub auto-backup every 2 hours (vs our 3am nightly)